Unsupervised prediction

Jeffrey Leek, Assistant Professor of Biostatistics
Johns Hopkins Bloomberg School of Public Health

Key ideas

- Sometimes you don't know the labels for prediction

- To build a predictor

    - Create clusters

    - Name clusters

    - Build predictor for clusters

- In a new data set

    - Predict clusters

# Iris example ignoring species labels

```
data(iris); library(ggplot2)
inTrain <- createDataPartition(y=iris$Species,
                               p=0.7, list=FALSE)
training <- iris[inTrain,]
testing <- iris[-inTrain,]
dim(training); dim(testing)
```

```
[1] 45  5
```

Cluster with k-means

```
kMeans1 <- kmeans(subset(training,select=-c(Species)),centers=3)
training$clusters <- as.factor(kMeans1$cluster)
qplot(Petal.Width,Petal.Length,colour=clusters,data=training)
```

Compare to real labels

```
table(kMeans1$cluster,training$Species)
```

```
    setosa versicolor virginica
1        0          1        23
2       35          0         0
3        0         34        12
```

Build predictor

```
modFit <- train(clusters ~.,data=subset(training,select=-c(Species)),method="rpart")
table(predict(modFit,training),training$Species)
```

```
    setosa versicolor virginica
  1      0          0        21
  2     35          0         0
  3      0         35        14
```

Apply on test

```
testClusterPred <- predict(modFit,testing)
table(testClusterPred ,testing$Species)
```

```
testClusterPred setosa versicolor virginica
              1      0          0        13
              2     15          0         0
              3      0         15         2
```

Notes and further reading

- The cl_predict function in the clue package provides similar functionality

- Beware over-interpretation of clusters!

- This is one basic approach to recommendation engines

- Elements of statistical learning

- Introduction to statistical learning